

OpenCV Package

opencv Package

- Open source computer vision package is a good source of image processing and computer vision algorithms.
- Several algorithms can be used to manipulate image and extract features that can be used later for machine learning approaches.
- Originally developed by Intel.
- The library is cross-platform and free for use under the open-source BSD license.
- OpenCV is written in C++ and its primary interface is in C++
- There are bindings in Python, Java, and MATLAB/OCTAVE
- Install opencv: `pip install opencv-python`

opencv Package

- Loading and saving images

import opencv

Read an image and
convert it into a gray
level image

Plot the image with
gray color map

Don't show axes

Use plt.show() to show the
image

get image dimensions

Save image with tif format

```
In [1]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import cv2
img = cv2.imread('noise_lung.png',cv2.IMREAD_GRAYSCALE)
plt.imshow(img,cmap='gray')
plt.axis('off')
plt.show()
```

Figure 1



```
In [2]: img.shape
```

```
Out[2]: (290, 400)
```

```
In [4]: cv2.imwrite('noise_lung.tif',img)
```

```
Out[4]: True
```

opencv Package

- Resizing and cropping images

Resize an image: provide the new dimensions as (width, height). You may specify the interpolation type as cubic interpolation



Use slicing to determine the region of the image you want to crop

Dimensions of the resized image (height, width)

Dimensions of the cropped image

```
In [13]: plt.figure(2)
plt.subplot(1,2,1)
im_2 = cv2.resize(img,(int(img.shape[1]*0.5),int(img.shape[0]*0.5)), interpolation=cv2.INTER_CUBIC)
plt.imshow(im_2,cmap='gray')
im_3 = img[73:218,100:300]
plt.subplot(1,2,2)
plt.imshow(im_3,cmap='gray')
plt.show()
```

Figure 2



resized

cropped

Forward to next view

```
In [9]: im_2.shape
out[9]: (145, 200)
```

```
In [12]: im_3.shape
out[12]: (145, 200)
```

What is image filtering?

$f(x,y)$



$g(x,y)$



filtering



filtering





Image Filtering Methods

- Spatial Domain




- Frequency Domain (i.e., uses Fourier Transform)



Area Shape and Size of the mask/kernel/filter

- Area shape is typically defined using a rectangular mask.
- Area size is determined by mask size.
e.g., 3x3 or 5x5
- Mask size is an important parameter!




Operation

- Typically, linear combinations of pixel values.
 - e.g., weight pixel values and add them together.
- Different results can be obtained using different weights.
 - (e.g., smoothing, sharpening, edge detection).

mask

w1	w2	w3
w4	w5	w6
w7	w8	w9



Example: Convolution Operation


1	-1	-1
-1	1	-1
-1	-1	1

Filter

1	0	0	0	0	0	1
0	1	0	0	0	1	0
0	0	1	0	0	0	0
1	0	0	0	0	1	0
0	1	0	0	0	1	0
0	0	1	0	0	1	0


6 x 6 image

Dot
product




Handling Pixels Close to Boundaries

pad with zeroes




wrap around



or

How do we choose the mask weights?

- Depends on the application.
- Usually by sampling certain functions and their derivatives.



Good for
image **smoothing**

Good for
image **sharpening**

Normalization of Mask Weights

- Sum of weights affects overall intensity of output image.
- Positive weights
 - Normalize them such that they sum to **one**.
- Both positive and negative weights
 - Should sum to **zero** (but not always)

w1	w2	w3
w4	w5	w6
w7	w8	w9

1	1	1	1	2	1
1/9	1	1	1	2	4
1	1	1	1	2	1

Smoothing Using Averaging

- **Idea:** replace each pixel by the average of its neighbors.
- Useful for reducing noise and unimportant details.
- The size of the mask controls the amount of smoothing.

$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

(a)

$$\frac{1}{25} \times \begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline \end{array}$$


(b)

$$\frac{1}{49} \times \begin{array}{|c|c|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline \end{array}$$

(c)

Smoothing Using Averaging (cont'd)


- **Trade-off:** noise vs blurring and loss of detail.



Gaussian Smoothing

- **Idea:** replace each pixel by a weighted average of its neighbors
- Mask weights are computed by sampling a Gaussian function

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}}$$





7×7 Gaussian mask


1	1	2	2	2	1	1
1	2	2	4	2	2	1
2	2	4	8	4	2	2
2	4	8	16	8	4	2
2	2	4	8	4	2	2
1	2	2	4	2	2	1
1	1	2	2	2	1	1

Note: weight values decrease with distance from mask center!


Gaussian Smoothing - Example



$\sigma = 1$ pixel



$\sigma = 5$ pixels



$\sigma = 10$ pixels



$\sigma = 30$ pixels

Averaging vs Gaussian Smoothing





Averaging



Gaussian

Image Sharpening

- Idea: compute intensity differences in local image regions.
- Useful for emphasizing transitions in intensity (e.g., in edge detection).




$$1/9 (-10 - 10 - 10 -10 + 80 -10 -10 -10 -10) = 0$$

(there is no variation in the gray-levels)


$$\rightarrow 1/9 (-10 - 80 - 80 -10 + 640 -80 -10 -80 -80) = 210/9 > 0$$

(there is variation in the gray-levels)

Example



before



after

opencv Package

- Image filtering:
Denoising/Smoothing

Use 5x5 Gaussian blurring/smoothing with $\sigma = 0.5$

Use 5x5 Gaussian blurring/smoothing with $\sigma = 4$

Use 5x5 median filter

Original


Gaussian Smoothing

Smoothing more as increasing σ

```
In [11]: im_denoise_1 = cv2.GaussianBlur(img,(5, 5), 0.5)
im_denoise_2 = cv2.GaussianBlur(img,(5, 5), 4)
im_denoise_3 = cv2.medianBlur(img, 5)

plt.figure(3)
plt.subplot(1,4,1)
plt.imshow(img,cmap='gray')
plt.title('original')
plt.axis('off')
plt.subplot(1,4,2)
plt.imshow(im_denoise_1,cmap='gray')
plt.title('denoised, '+'$\sigma = 0.5$')
plt.axis('off')
plt.subplot(1,4,3)
plt.imshow(im_denoise_2,cmap='gray')
plt.title('denoised, '+'$\sigma = 4$')
plt.axis('off')
plt.subplot(1,4,4)
plt.imshow(im_denoise_3,cmap='gray')
plt.title('denoised, median')
plt.axis('off')
plt.show()
```

Figure 3



Median filtering

opencv Package

- Image filtering: sharpening, and edge preserve filters

Define 3x3
sharpening filter

Apply the filter

Apply bilateral
filter to denoise and
preserve edges from
blurring


Original

Gaussian Smoothing

Sharpening after smoothing

```
In [19]: kernel = np.array([[-1,-1,-1],  
                         [-1, 9,-1],  
                         [-1,-1,-1]])  
  
im_sharp = cv2.filter2D(im_denoise_2,-1, kernel)  
plt.figure(4)  
plt.subplot(1,4,1)  
plt.imshow(img,cmap='gray')  
plt.title('original')  
plt.axis('off')  
plt.subplot(1,4,2)  
plt.imshow(im_denoise_2,cmap='gray')  
plt.title('denoised, '+r'$\sigma = 4$')  
plt.axis('off')  
plt.subplot(1,4,3)  
plt.imshow(im_sharp,cmap='gray')  
plt.title('sharpening')  
plt.axis('off')  
plt.subplot(1,4,4)  
im_edge_preserve = cv2.bilateralFilter(img,9,31,31)  
plt.imshow(im_edge_preserve,cmap='gray')  
plt.title('denoised, edge_preserve')  
plt.axis('off')  
plt.show()
```

Figure 4



Smoothing with edge preserve filtering

opencv Package

- Image binarization using adaptive thresholding
the output pixel either zero or max_gray

Sliding window size e.g.
301x301

Shift threshold by this value to fine tune the output


- There are two options: cv2.THRESH_BINARY or cv2.THRESH_BINARY_INV
You choose a binary output or inverted binary
- The thresholding approach is either Gaussian (weighted average) or mean.

Input image Thresholding approach

```
max_gray = 255
local_window_size = 301
delta_threshold = 0
im_binary_1 = cv2.adaptiveThreshold(im_edge_preserve, max_gray, cv2.ADAPTIVE_THRESH_GAUSSIAN_C,
                                    cv2.THRESH_BINARY, local_window_size, delta_threshold)
im_binary_2 = cv2.adaptiveThreshold(im_edge_preserve, max_gray, cv2.ADAPTIVE_THRESH_MEAN_C,
                                    cv2.THRESH_BINARY, local_window_size, delta_threshold)

plt.figure(5)
plt.subplot(1,2,1)
plt.imshow(im_binary_1,cmap='gray')
plt.title('adaptive binarization (Gaussian)')
plt.axis('off')
plt.subplot(1,2,2)
plt.imshow(im_binary_2,cmap='gray')
plt.title('adaptive binarization (mean)')
plt.axis('off')
plt.show()
```

Figure 5



cv2.ADAPTIVE_THRESH_MEAN_C: The threshold value is the mean of the neighbourhood area minus the constant C.

cv2.ADAPTIVE_THRESH_GAUSSIAN_C: The threshold value is a gaussian-weighted sum of the neighbourhood values minus the constant C

opencv Package

- **Image binarization using adaptive thresholding**



opencv Package

- Edge Detection: There are several algorithms such as:

- Sobel,
- Laplacian,
- Canny

Laplacian edge detector


Canny edge detector,
Lower threshold = 20
Upper threshold = 30
Stronger edges are above
the upper and weaker
edges are greater than a
lower threshold and less
than the upper one.

Input image (binary image)


Output image bit depth

```
In [49]:  
plt.figure(6)  
plt.subplot(1,2,1)  
im_edge_1 = cv2.Laplacian(im_binary_2,cv2.CV_8U)  
plt.imshow(im_edge_1,cmap='gray')  
plt.title('Laplacian Edge Detector')  
plt.axis('off')  
plt.subplot(1,2,2)  
im_edge_2 = cv2.Canny(im_binary_2,20,30)  
plt.imshow(im_edge_2,cmap='gray')  
plt.title('Canny Edge Detector')  
plt.axis('off')  
plt.show()
```


Figure 6



Laplacian



Canny



opencv Package

- Corner Detection: There are several algorithms such as Harris, minimum eigenvalue and FAST

Input image

Find the cornerness measure

It is not needed for the algorithm. It is just to replicate corners to look in a good size to visualize

Threshold to find the strongest corners. Then overlay these corners (white pixels) on top of the original image

Window size to calculate the cornerness

Free parameter k used in calculating the cornerness

Filter size to calculate the gradient

The screenshot shows a Jupyter Notebook cell with the following Python code:

```
cornerness = cv2.cornerHarris(np.float32(im_denoise_2), 3, 3, 0.04)
cornerness = cv2.dilate(cornerness, None)
overlay_corners = np.copy(im_denoise_2)
overlay_corners[cornerness > 0.01 * cornerness.max()] = 255
plt.figure(7)
plt.subplot(1,2,1)
plt.imshow(im_denoise_2, cmap='gray')
plt.title('input')
plt.axis('off')
plt.subplot(1,2,2)
plt.imshow(overlay_corners, cmap='gray')
plt.title('detected corners')
plt.axis('off')
```

Below the code, there are two images labeled "input" and "detected corners". The "input" image is a grayscale scan of lungs. The "detected corners" image shows the same lungs with white dots overlaid at the detected corner points.

opencv Package

- Processing Color Images



Read a color image

Convert image from
BGR to the regular
RGB. (opencv uses
BGR order by default)

Matplotlib plots the
image as RGB while
opencv read it as BGR.
that is why the R and B
are swapped in the plot

```
In [1]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import cv2
img = cv2.imread('fundus_1.png',cv2.IMREAD_COLOR)
plt.figure(1)
plt.subplot(1,2,1)
plt.axis('off')
plt.imshow(img)
img1 = cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.subplot(1,2,2)
plt.axis('off')
plt.imshow(img1)
plt.show()
```

Figure 1



Correct RGB order

Reset original view

opencv Package


- Color Image Components

Color image three components

Plot the three components separately

```
In [2]: r = img1[:, :, 0]
g = img1[:, :, 1]
b = img1[:, :, 2]
plt.figure(2)
plt.subplot(1,3,1), plt.axis('off')
plt.imshow(r,cmap='gray')
plt.subplot(1,3,2), plt.axis('off')
plt.imshow(g,cmap='gray')
plt.subplot(1,3,3), plt.axis('off')
plt.imshow(b,cmap='gray')
plt.show()
```

Figure 2



Red

Green

Blue



Download plot

opencv Package


- **Image Histogram**

Use `calcHist()` to find
the histogram of red
with 256 bins

Histogram of Green

Histogram of Blue

Plot of R, G, and B
histograms



opencv Package


- Histogram Equalization

Convert image to YUV color space, where Y is the intensity and U and V are chroma components.

Histogram of Y before equalization

Apply equalization to Y

input



equalized

```
In [4]:  
img_yuv = cv2.cvtColor(img1, cv2.COLOR_RGB2YUV)  
hist_y_before = cv2.calcHist([img_yuv], [0], None, [256], [0,256])  
img_yuv[:, :, 0] = cv2.equalizeHist(img_yuv[:, :, 0])  
hist_y_after = cv2.calcHist([img_yuv], [0], None, [256], [0,256])  
img_rgb = cv2.cvtColor(img_yuv, cv2.COLOR_YUV2RGB)  
plt.figure(4)  
plt.subplot(1, 2, 1)  
plt.imshow(img1), plt.axis('off')  
plt.subplot(1, 2, 2)  
plt.imshow(img_rgb), plt.axis('off')  
plt.show()
```

Figure 4

Histogram of Y
After equalization

opencv Package


- Histogram Equalization

Plot histograms



```
In [5]: plt.figure(5)
plt.subplot(1,2,1)
plt.plot(hist_y_before, 'r')
plt.xlim([0,256])
plt.subplot(1,2,2)
plt.plot(hist_y_after, 'g')
plt.xlim([0,256])
plt.show()
```

Y histogram before
equalization (levels are not
uniformly distributed)



You need to convert RGB to
BGR to write the image in the
standard RGB order



```
In [7]: cv2.imwrite('img_before_eq.png', cv2.cvtColor(img1, cv2.COLOR_RGB2BGR))
cv2.imwrite('img_after_eq.png', cv2.cvtColor(img_rgb, cv2.COLOR_RGB2BGR))
```

Out[7]: True

Y histogram after equalization equalized, more uniform, more contrast

